

Class 7: Machine learning 2


Assael Madrigal (PID: A10179083)

#Clustering k-means clustering is very prevalent. k means that we need to tell a k - how many groups I want, later we can tell it what it should be after analysing the output but we have to start with something.

To get started let's make some data, let's see how rnorm works by plotting a histogram

```
#rnorm generates as many random numbers as I ask drawn from a normal distribution  
  
hist(rnorm(10000, mean=3))
```

Histogram of rnorm(10000, mean = 3)



```
#so the mean is where the middle of the histogram is going to be
```


here we are going to make 2 groupings one centered around 3 and the other at -30

```
tmp <- c(rnorm(30, mean=3), rnorm(30,-3))
x <- cbind(x=tmp, y=rev(tmp))
x
```

	x	y
[1,]	3.88368344	-4.88962659
[2,]	4.50401670	-3.20608635
[3,]	2.17771563	-3.17241391
[4,]	2.30845033	-3.94125846
[5,]	1.87568713	-3.08629966
[6,]	3.09716803	-3.48682628
[7,]	1.66885452	-3.60948746
[8,]	3.74973276	-4.83342587
[9,]	1.13218560	-2.91836322
[10,]	2.26618438	-3.03263986
[11,]	2.71943019	-3.63285281
[12,]	1.41917840	-4.04094533
[13,]	2.99648374	-3.60736534
[14,]	2.60611441	-2.35027179
[15,]	4.06415916	-2.88471970
[16,]	2.81152556	-1.17655573
[17,]	5.68402468	-1.79511979
[18,]	1.69687925	-2.07273771
[19,]	1.68647329	-3.80172581
[20,]	4.67522726	-3.73724432
[21,]	2.28062900	-0.91999295
[22,]	2.93179458	-3.44172582
[23,]	3.78364966	-1.74836331
[24,]	3.43217575	-4.19606067
[25,]	3.21032670	-2.57716698
[26,]	3.21343697	0.07147265
[27,]	2.50009830	-3.11323793
[28,]	3.68938747	-3.19025763
[29,]	3.03551624	-0.49142900
[30,]	2.86943934	-4.44232525
[31,]	-4.44232525	2.86943934
[32,]	-0.49142900	3.03551624
[33,]	-3.19025763	3.68938747
[34,]	-3.11323793	2.50009830
[35,]	0.07147265	3.21343697

[36,]	-2.57716698	3.21032670
[37,]	-4.19606067	3.43217575
[38,]	-1.74836331	3.78364966
[39,]	-3.44172582	2.93179458
[40,]	-0.91999295	2.28062900
[41,]	-3.73724432	4.67522726
[42,]	-3.80172581	1.68647329
[43,]	-2.07273771	1.69687925
[44,]	-1.79511979	5.68402468
[45,]	-1.17655573	2.81152556
[46,]	-2.88471970	4.06415916
[47,]	-2.35027179	2.60611441
[48,]	-3.60736534	2.99648374
[49,]	-4.04094533	1.41917840
[50,]	-3.63285281	2.71943019
[51,]	-3.03263986	2.26618438
[52,]	-2.91836322	1.13218560
[53,]	-4.83342587	3.74973276
[54,]	-3.60948746	1.66885452
[55,]	-3.48682628	3.09716803
[56,]	-3.08629966	1.87568713
[57,]	-3.94125846	2.30845033
[58,]	-3.17241391	2.17771563
[59,]	-3.20608635	4.50401670
[60,]	-4.88962659	3.88368344

```
plot(x)
```



The main function in R for K-means clustering is called `kmeans()`, `kmeans(x, centers,...)`. The center is the number of clusters. `nstart` is the number of iterations `kmeans` will go, so one way is to keep increasing it until the answer does not change or! even better you can plot the “scree plot” and look at the elbow.

But it is still a limitation, but the advantage is that `kmeans` is very fast.

```
k <- kmeans(x,centers=2, nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.977502	2.932321
2	2.932321	-2.977502

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 74.34935 74.34935
(between_SS / total_SS = 87.6 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q1. How many points do I need to cluster? (from k)

```
k$size
```

```
[1] 30 30
```

```
#a vector with the sizes of each cluster
```

Q2. The clustering result ie membership vector?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q3. Cluster centers?

```
k$centers
```

```
      x      y
1 -2.977502  2.932321
2  2.932321 -2.977502
```


Q4. Make a plot of our data colored by clustering results with optionally the clusters centers shown

```
#col is color but if we give it an number it has a color assigned to it, so for this case
plot(x, col=k$cluster, pch=16)
points(k$centers, col="blue", pch=15, cex=2)
```



Q.5 Run kmeans again but cluster into 3 groups and plot the results

```
k3 <- kmeans(x,centers=3, nstart = 20)
plot(x, col=k3$cluster, pch=16)
points(k3$centers, col="blue", pch=15, cex=2)
```



The main problem with kmeans is that it will fit the data into the structure, so if we give it 3 clusters it will split it into 3 clusters. So on, so we have to be careful.

kmeans: - breaks observations into k-predefined number of clusters - you define the number of clusters! - help by plotting with scree plots

Heirarchical clustering

Heirarchical clustering has the advantage that it can potentially reveal structure in the data rather than imposing one as k-means will.

The main function in “base R” is `hclust()`, is follows `hclust(d, method = "complete", members = NULL)` where ‘d’ was produced by `dist()` or any measure of dissimilarity.

It requires a distance matrix as input, not the raw data itself

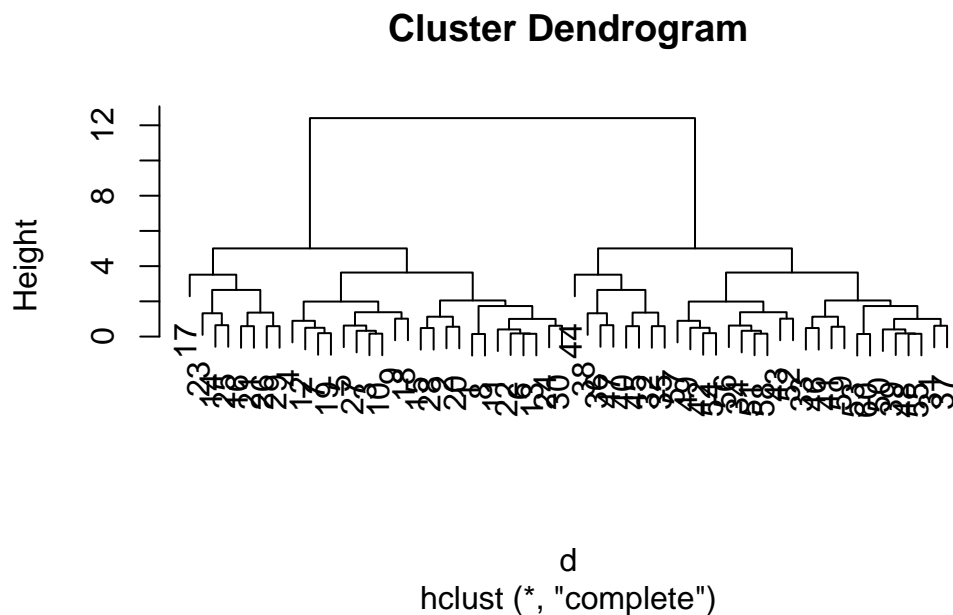
```
d<-dist(x)
hc <- hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

```
plot(hc)
```



The crossbar height is how far apart the datapoints are. So in this case the largest difference is in the first 2 clusters and then the smaller clusters are very little apart from there. Y axis, height, is the distance of the two branches below it.

There are two forms of hclust: bottom up vs top down

The function to get our clusters/groups from a `hclust` object is called `cutree()`

```
cutree(hc, k=2)
```

[illegible]



```
#or
cutree(hc, h=8)
```

[1] 1 2 2 2 2 2 2 2 2
[39] 2


```
grps <- cutree(hc, k=2)
```

Q. plot our hclust results in terms of our data colored by cluster membership

```
plot(x, col=grps)
```



```
plot(x, col=grps, pch=15) #pch is just the shape of the points
```



Principal Component Analysis (PCA)

UK food class lab

first import the data

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
## Complete the following code to find out how many rows and columns are in x?  
str(x)
```

```
'data.frame':  17 obs. of  5 variables:
```

```
$ X      : chr  "Cheese" "Carcass_meat " "Other_meat " "Fish" ...
$ England : int  105 245 685 147 193 156 720 253 488 198 ...
$ Wales   : int  103 227 803 160 235 175 874 265 570 203 ...
$ Scotland : int  103 242 750 122 184 147 566 171 418 220 ...
$ N.Ireland: int   66 267 586 93 209 139 1033 143 355 187 ...
```

```
dim(x)
```

```
[1] 17  5
```

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

17 rows (observations) and 5 columns (variables) it is using the names of the food as a column so we are going to make it a name instead

```
rownames(x)
```

```
[1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10" "11" "12" "13" "14" "15"
[16] "16" "17"
```

I can change them like this: but this is overwriting x every time, because it is removing a column every time.

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267

Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

a better way to do this is: when we read the data we can tell it that the row names is in column 1 `row.names=1`

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer assigning the row names at the same time as reading the data in. The problem may be if there are no rownames, but I would get the data first then visualize with `head()` and then decide which column to pick for rownames

Q3: Changing what optional argument in the above `barplot()` function results in the following plot? ()

I just set `beside` to `False`

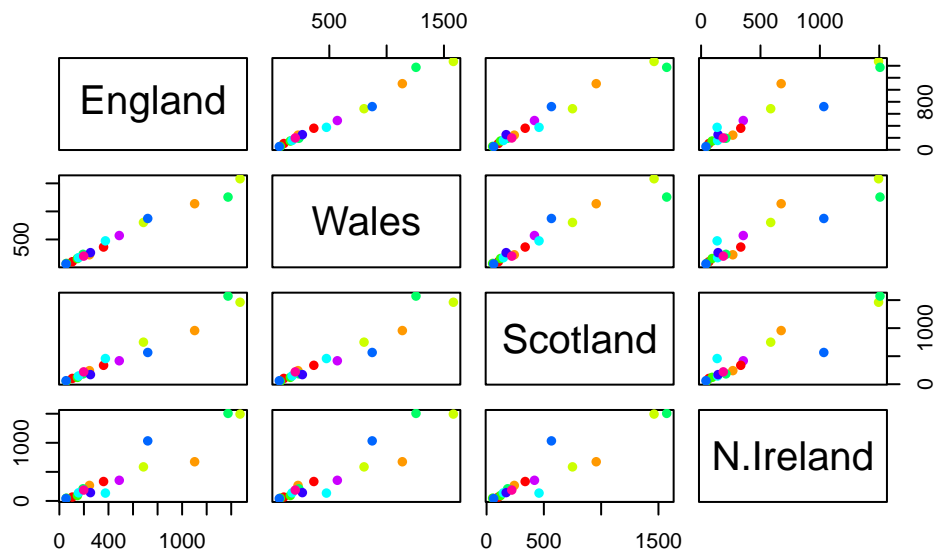
```
barplot(as.matrix(x), beside=F,col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

So the diagonal means they consume the same amount for that particular food. Deviation from the diagonal is what is different between them. Each graph is comparing two countries to each other, but it is mirrored. The main thing is that the difference is what we are looking for.

```
pairs(x, col=rainbow(10), pch=16)
```



#Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Based on the graphs alone it is the food represented by the blue dot, the orange dot. but it is really hard to tell.

To help me makes sense of this data... The main function for PCA in base R is called `prcomp()`

It wants the transpose (with the `t()`) of our food data for analysis

```
pca <- prcomp(t(x))

summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

The proportion of variance is what i am looking at: 0.67 for PC1 it means that PC1 is responsible for 67% of the variance. The cumulative variance is the sum of them so for PC1 and PC2 it is 96.5%

One of the main results that folks look for is called the “score plot” a.k.a. PC plot, PC1 vs PC2 plot...

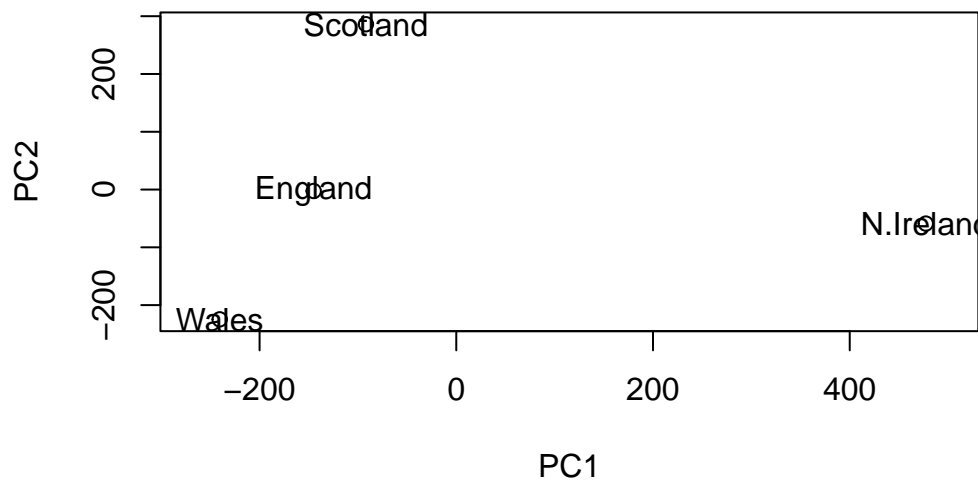
We are looking at the distance from the point to the 0,0

```
plot(pca$x[,1], pca$x[,2])  
abline(h=0, v=0, col="gray", lty=2) #adding a line just to see where 0,0 is
```




Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2  
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))  
text(pca$x[,1], pca$x[,2], colnames(x))
```



#Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange","red","blue","green"))
```

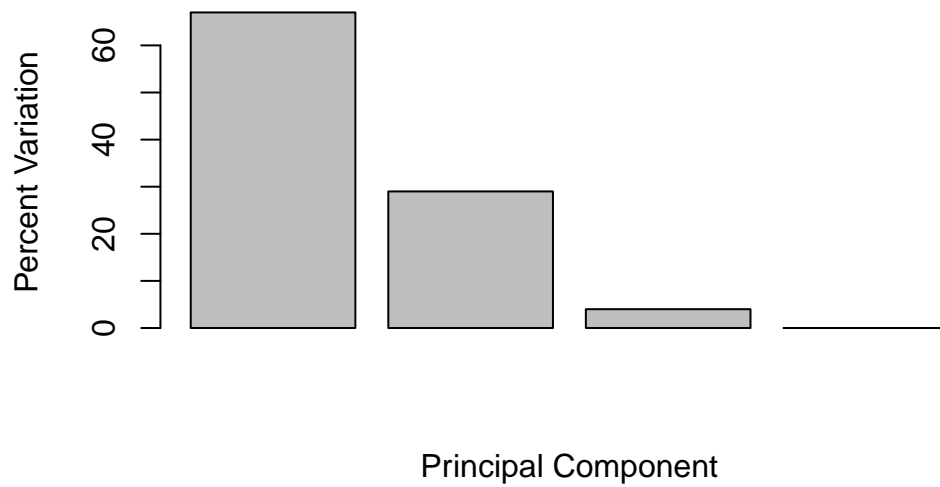
```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```

```
## or the second row here...
z <- summary(pca)
z$importance
```

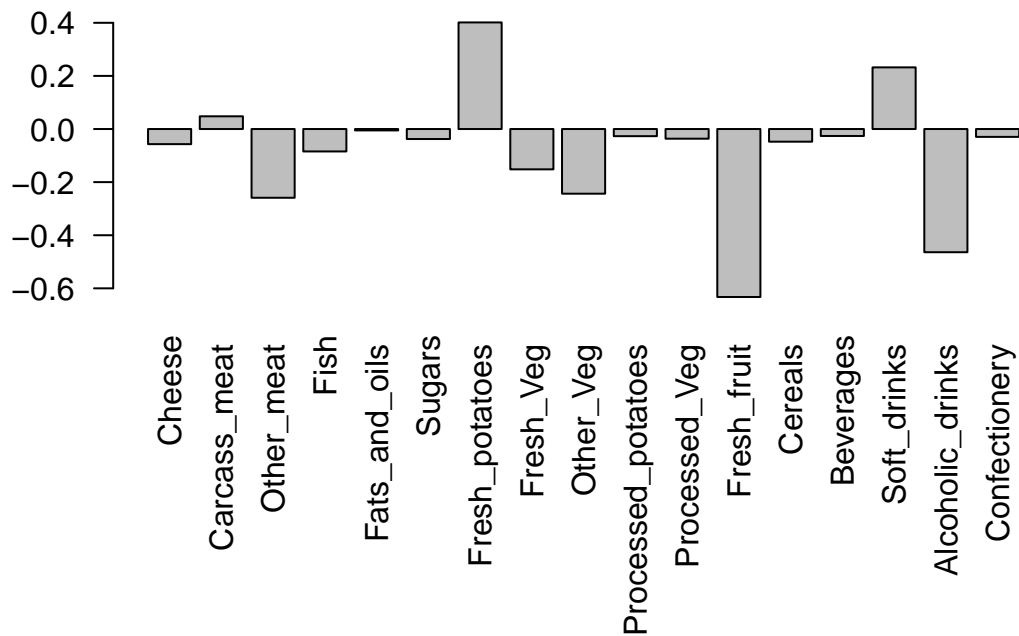
	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	3.175833e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
#This information can be summarized in a plot of the variances (eigenvalues) with respect
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



To dig deeper we can see how much each variable affects the original PCA1 from the `$rotation` from `prcomp()` and summarized in `biplot()`

```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```

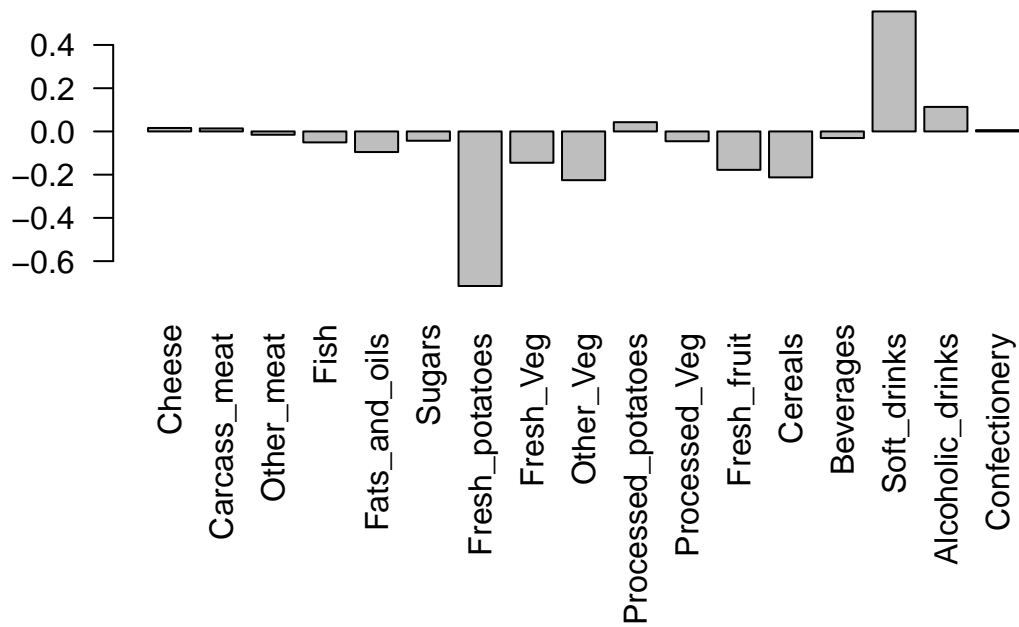


The two biggest contributors are fresh fruit and soft drinks because they are the biggest bars.

Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

Fresh potatoes makes the countries go left and soft drinks right so fresh potatoes differ the most between N. Ireland and Soft drinkg makes it the most different.

```
#PC2
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```




PC2 is predominantly still soft drinks but also alcoholic drinks.


```
library(ggplot2)

df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")

# Our first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```




```
ggplot(df_lab) +  
  aes(PC1, PC2, col=Country, label=Country) +  
  geom_hline(yintercept = 0, col="gray") +  
  geom_vline(xintercept = 0, col="gray") +  
  geom_point(show.legend = FALSE) +  
  geom_label(hjust=1, nudge_x = -10, show.legend = FALSE) +  
  expand_limits(x = c(-300,500)) +  
  xlab("PC1 (67.4%)") +  
  ylab("PC2 (28%)") +  
  theme_bw()
```



we can also make a nice ggplot for our pca


```
ld <- as.data.frame(pca$rotation)
ld_lab <- tibble::rownames_to_column(ld, "Food")

ggplot(ld_lab) +
  aes(PC1, Food) +
  geom_col()
```




and we can also make it nicer to see by color scaling and ordering them by highest to smallest

```
ggplot(ld_lab) +
  aes(PC1, reorder(Food, PC1), bg=PC1) +
  geom_col() +
  xlab("PC1 Loadings/Contributions") +
  ylab("Food Group") +
  scale_fill_gradient2(low="purple", mid="gray", high="darkgreen", guide=NULL) +
  theme_bw()
```



another way to visualize this is by using a biplot

```
## The inbuilt biplot() can be useful for small datasets  
biplot(pca)
```

PCA for RNA-seq

loading the data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set?

```
class(rna.data)
```

```
[1] "data.frame"
```

```
str(rna.data)
```

```
'data.frame':  100 obs. of  10 variables:
 $ wt1: int  439 219 1006 783 181 460 27 175 658 121 ...
 $ wt2: int  458 200 989 792 249 502 30 182 669 116 ...
 $ wt3: int  408 204 1030 829 204 491 37 184 653 134 ...
 $ wt4: int  429 210 1017 856 244 491 29 166 633 117 ...
 $ wt5: int  420 187 973 760 225 493 34 180 657 133 ...
 $ ko1: int   90 427 252 849 277 612 304 255 628 931 ...
 $ ko2: int   88 423 237 856 305 594 304 291 627 941 ...
 $ ko3: int   86 434 238 835 272 577 285 305 603 990 ...
 $ ko4: int   90 433 226 885 270 618 311 271 635 982 ...
 $ ko5: int   93 426 210 894 279 638 285 269 620 934 ...
```

```
dim(rna.data)
```

```
[1] 100  10
```

There are 100 genes and 10 samples (columns)